

In [57]:

```
import pandas as pd
import numpy as np
import os
import matplotlib.pyplot as plt
import seaborn as sns
```

In [58]:

```
#Loading Dataset
df_trvl_rvw = pd.read_csv('tripadvisor_review.csv')
#df_trvl_rvw = df_trvl_rvw.set_index('User ID')
df_trvl_rvw.head()
```

Out[58]:

	User ID	Category 1	Category 2	Category 3	Category 4	Category 5	Category 6	Category 7	Category 8	Cε
0	User 1	0.93	1.8	2.29	0.62	0.80	2.42	3.19	2.79	
1	User 2	1.02	2.2	2.66	0.64	1.42	3.18	3.21	2.63	
2	User 3	1.22	0.8	0.54	0.53	0.24	1.54	3.18	2.80	
3	User 4	0.45	1.8	0.29	0.57	0.46	1.52	3.18	2.96	
4	User 5	0.51	1.2	1.18	0.57	1.54	2.02	3.18	2.78	

In [59]:

```
#rename the columns for easier understanding
re_name = ['User ID','art galleries',
           'dance clubs',
           'juice bars',
           'restaurants',
           'museums',
           'resorts',
           'parks',
           'beaches',
           'theaters',
           'religious']

#rename
df_trvl_rvw.columns = re_name
df_trvl_rvw.head()
```

Out[59]:

	User ID	art galleries	dance clubs	juice bars	restaurants	museums	resorts	parks	beaches	theaters	religious
0	User 1	0.93	1.8	2.29	0.62	0.80	2.42	3.19	2.79	1.82	
1	User 2	1.02	2.2	2.66	0.64	1.42	3.18	3.21	2.63	1.86	
2	User 3	1.22	0.8	0.54	0.53	0.24	1.54	3.18	2.80	1.31	
3	User 4	0.45	1.8	0.29	0.57	0.46	1.52	3.18	2.96	1.57	
4	User 5	0.51	1.2	1.18	0.57	1.54	2.02	3.18	2.78	1.18	

In [60]:

```
df_trvl_rvw.isnull().sum()
```

Out[60]:

User ID	0
art galleries	0
dance clubs	0
juice bars	0
restaurants	0
museums	0
resorts	0
parks	0
beaches	0
theaters	0
religious	0
	dtype: int64

In [61]:

```
df_trvl_rvw.drop('User ID', axis = 1, inplace = True)
df_trvl_rvw = df_trvl_rvw.fillna(0)
```

In [62]:

```
df_trvl_rvw.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 980 entries, 0 to 979
Data columns (total 10 columns):
 #   Column           Dtype  
 --- 
 0   art galleries    float64
 1   dance clubs      float64
 2   juice bars       float64
 3   restaurants      float64
 4   museums          float64
 5   resorts          float64
 6   parks            float64
 7   beaches          float64
 8   theaters          float64
 9   religious         float64
dtypes: float64(10)
memory usage: 76.7 KB
```

In [63]:

```
#All of our data is type 'float'
pd.set_option('display.max_columns', 30)
df_trvl_rvw.describe()
```

Out[63]:

	art galleries	dance clubs	juice bars	restaurants	museums	resorts	parks	
<b>count</b>	980.000000	980.000000	980.000000	980.000000	980.000000	980.000000	980.000000	98
<b>mean</b>	0.893194	1.352612	1.013306	0.532500	0.939735	1.842898	3.180939	
<b>std</b>	0.326912	0.478280	0.788607	0.279731	0.437430	0.539538	0.007824	
<b>min</b>	0.340000	0.000000	0.130000	0.150000	0.060000	0.140000	3.160000	
<b>25%</b>	0.670000	1.080000	0.270000	0.410000	0.640000	1.460000	3.180000	
<b>50%</b>	0.830000	1.280000	0.820000	0.500000	0.900000	1.800000	3.180000	
<b>75%</b>	1.020000	1.560000	1.572500	0.580000	1.200000	2.200000	3.180000	
<b>max</b>	3.220000	3.640000	3.620000	3.440000	3.300000	3.760000	3.210000	

In [64]:

```
# visualize our first plot: we will examine the number of reviews under each category
import matplotlib.pyplot as plt
import warnings
%matplotlib inline
import plotly.express as px

# Plotting pretty figures and avoid blurry images
%config InlineBackend.figure_format = 'retina'
# Larger scale for plots in notebooks

# Ignore warnings
import warnings
warnings.filterwarnings('ignore')

# Enable multiple cell outputs
from IPython.core.interactiveshell import InteractiveShell
InteractiveShell.ast_node_interactivity = 'all'

column_names = ['art galleries',
'dance clubs',
'juice bars',
'restaurants',
'museums',
'resorts',
'parks',
'beaches',
'theaters',
'religious']

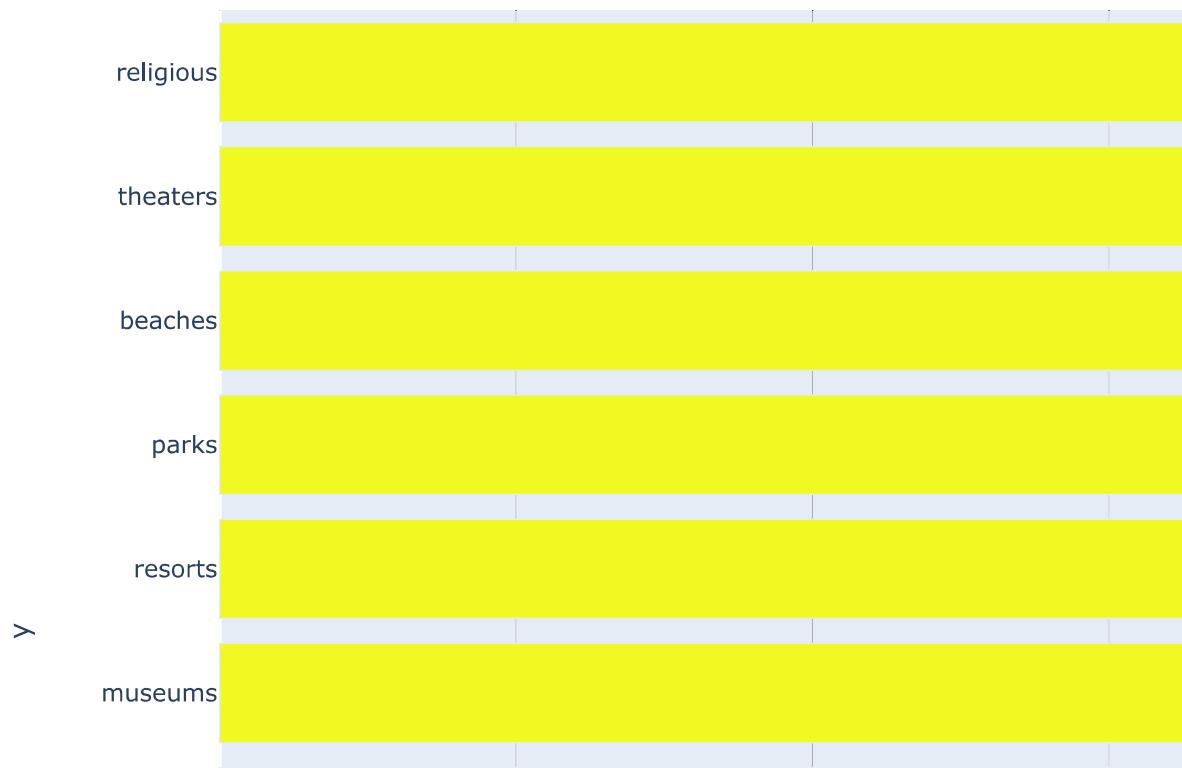
counts = df_trvl_rvw[column_names[:]].astype(bool).sum(axis=0).sort_values()

test = []
for i in range(len(counts.index)):
    test.append(counts.index[i])

fig = px.bar(counts,
            x=counts,
            y=test,
            color=counts,
            labels={
                "Total Ratings": "this is x",
                "Categories": "this is y"
            },
            height = 800,
            title="Number of reviews under each category")

fig.show()
```

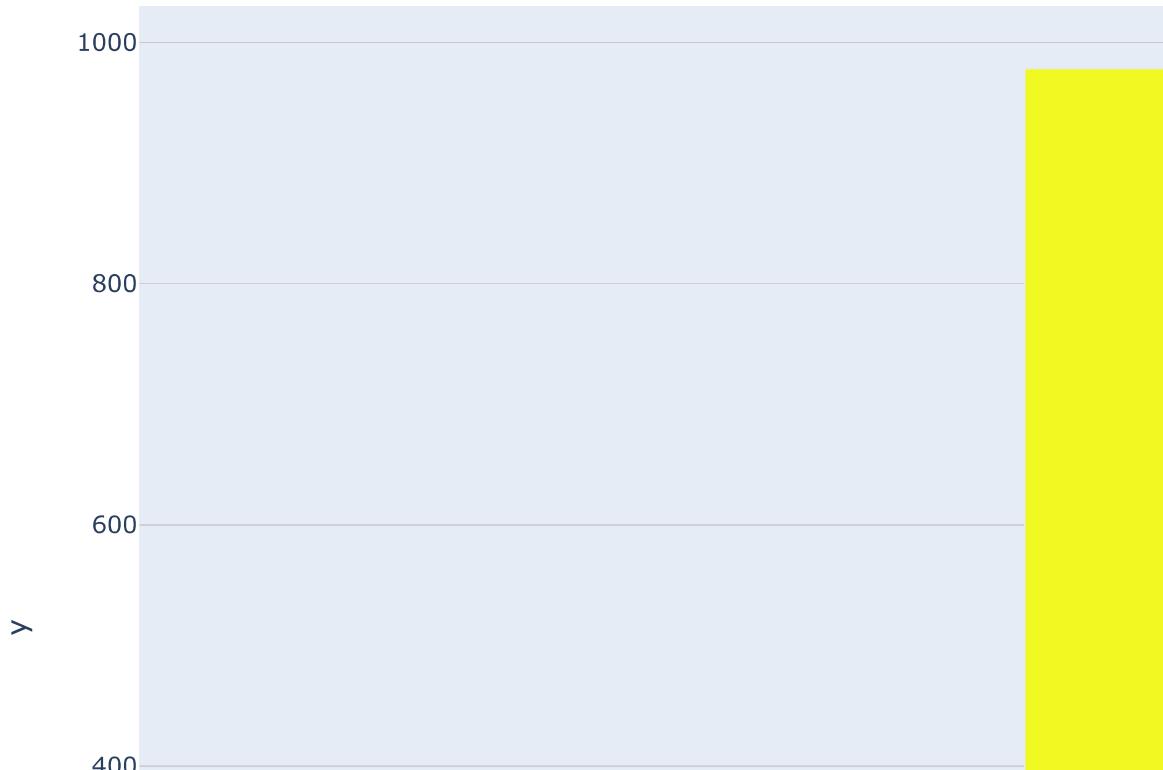
Number of reviews under each category



In [65]:

```
#reviews were given for each category.  
reviews = df_trvl_rvw[column_names[:]].astype(bool).sum(axis=1).value_counts()  
  
fig = px.bar(reviews,  
              x=reviews.index,  
              y=reviews.values,  
              color=reviews.values,  
              height = 800,  
              title="Number of categories VS Number of reviews")  
  
fig.show()
```

Number of categories VS Number of reviews





In [66]:

```
#average rating for each category
avg_rating = df_trvl_rvw[column_names[:]].mean().sort_values()

fig = px.bar(avg_rating,
              x = avg_rating.index,
              y = avg_rating.values,
              color = avg_rating.values,
              height = 800,
              title = "Average rating for each category")

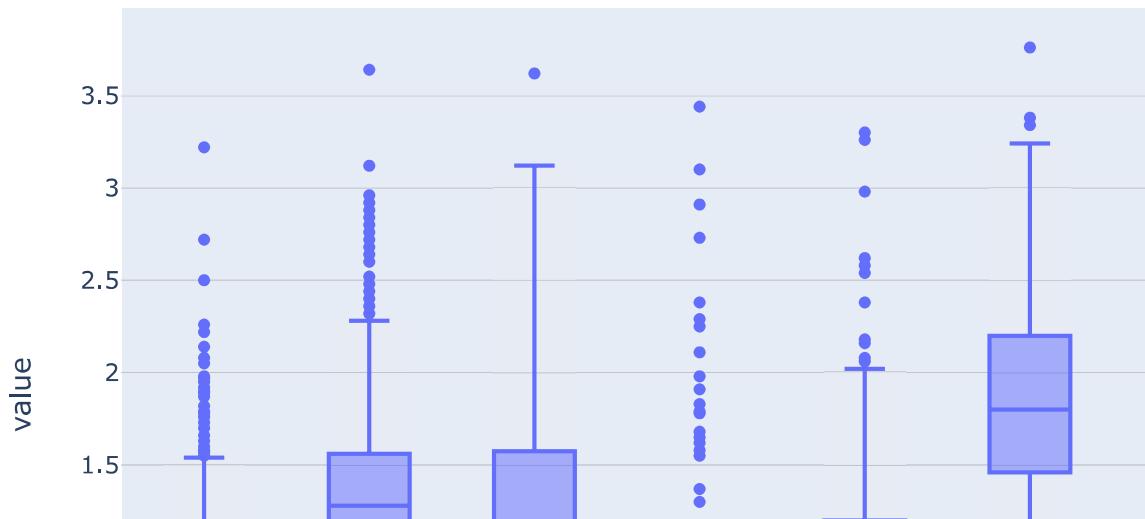
fig.show()
```

Average rating for each category



In [67]:

```
fig = px.box(df_trvl_rvw, y = ['art galleries'  
, 'dance clubs'  
, 'juice bars'  
, 'restaurants'  
, 'museums'  
, 'resorts'  
, 'parks'  
, 'beaches'  
, 'theaters'  
, 'religious'])  
fig.show()
```



In [68]:

```
df_new = pd.DataFrame(df_trvl_rvw)

Q1 = df_new.quantile(0.25)
Q3 = df_new.quantile(0.75)
IQR = Q3 - Q1

pd.set_option('display.max_info_rows', 30)

((df_new < (Q1 - 1.5 * IQR)) | (df_new > (Q3 + 1.5 * IQR))).sum()
```

Out[68]:

```
art galleries    47
dance clubs     54
juice bars      1
restaurants     45
museums         13
resorts          4
parks           399
beaches          24
theaters         23
religious        0
dtype: int64
```

In [69]:

```
import seaborn as sns

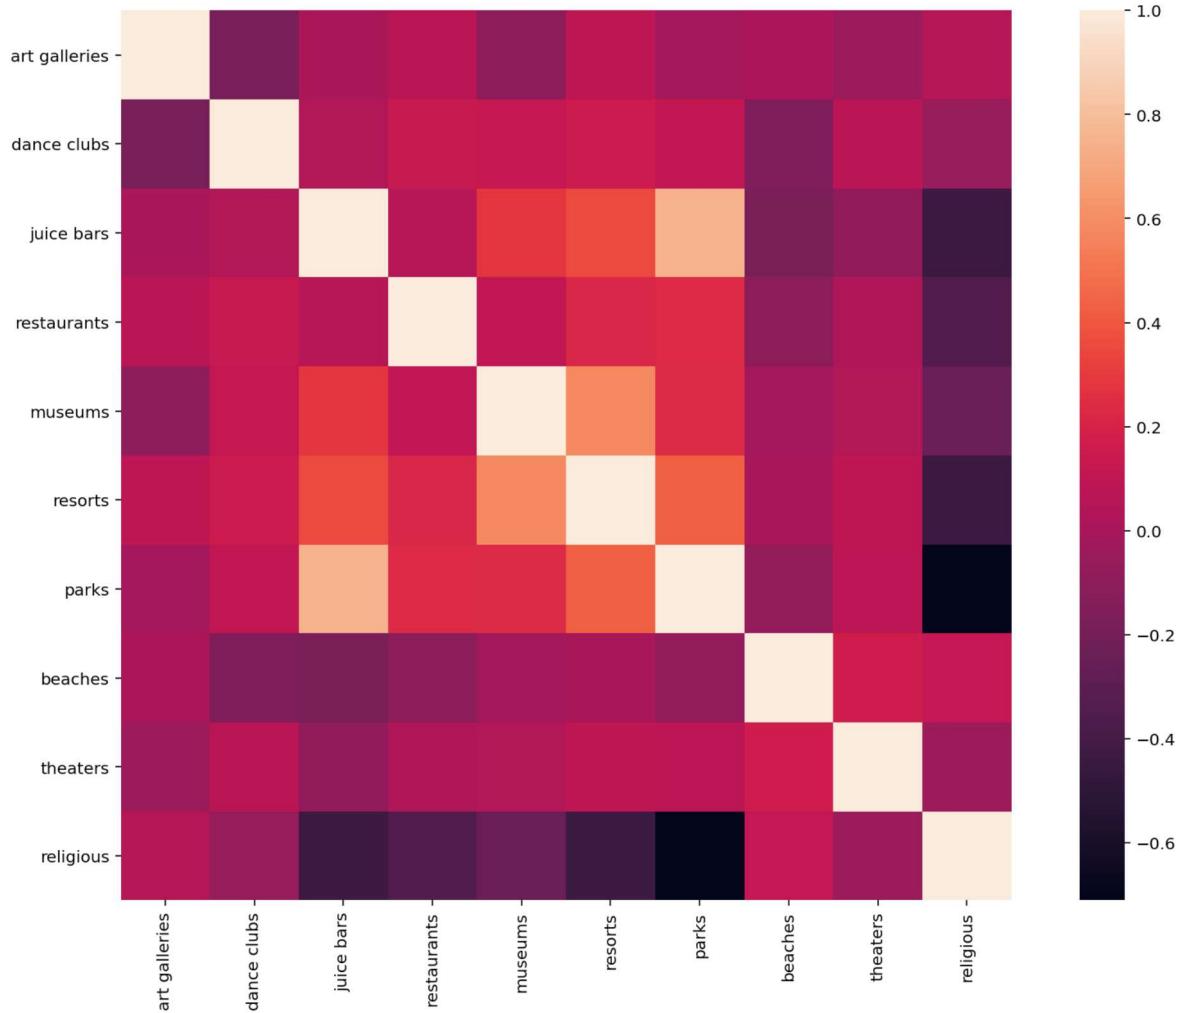
plt.figure(figsize=(15,10))
cor = df_trvl_rvw.corr() #Calculate the correlation of the above variables
sns.heatmap(cor, square = True) #Plot the correlation as heat map
```

Out[69]:

&lt;Figure size 1080x720 with 0 Axes&gt;

Out[69]:

&lt;AxesSubplot:&gt;



In [70]:

```
df_trvl_rvw.shape
```

Out[70]:

```
(980, 10)
```

In [71]:

```
df_trvl_rvw = df_trvl_rvw.drop_duplicates()
```

In [72]:

```
df_trvl_rvw.shape
```

Out[72]:

```
(941, 10)
```

In [73]:

```
entertainment = ['theaters', 'dance clubs']
food_travel = ['restaurants', 'resorts', 'juice bars']
historical = ['museums', 'art galleries', 'religious']
nature = ['beaches', 'parks']
```

In [74]:

```
df_categories = pd.DataFrame(columns = ['entertainment', 'food_travel', 'historical', 'natu
```

In [75]:

```
df_categories['entertainment'] = df_trvl_rvw[entertainment].mean(axis=1)
df_categories['food_travel'] = df_trvl_rvw[food_travel].mean(axis = 1)
df_categories['historical'] = df_trvl_rvw[historical].mean(axis = 1)
df_categories['nature'] = df_trvl_rvw[nature].mean(axis = 1)
```

In [76]:

```
df_categories.describe()
```

Out[76]:

	entertainment	food_travel	historical	nature
<b>count</b>	941.000000	941.000000	941.000000	941.000000
<b>mean</b>	1.462912	1.121647	1.546213	3.008799
<b>std</b>	0.312203	0.388452	0.186365	0.068679
<b>min</b>	0.500000	0.156667	1.106667	2.810000
<b>25%</b>	1.245000	0.820000	1.413333	2.960000
<b>50%</b>	1.445000	1.060000	1.530000	3.005000
<b>75%</b>	1.625000	1.410000	1.656667	3.050000
<b>max</b>	2.775000	2.773333	2.336667	3.285000

In [77]:

```
fig = px.box(df_categories, y=['entertainment', 'food_travel', 'historical', 'nature'])
fig.show()
```



In [78]:

```
from pyclustertend import hopkins  
  
hopkins(df_categories, df_categories.shape[0])
```

Out[78]:

0.21364958744784734

In [79]:

```
#Agglomerative Clustering
from scipy.cluster.hierarchy import dendrogram
from scipy.cluster import hierarchy
from sklearn.cluster import AgglomerativeClustering
from scipy.spatial.distance import pdist
import plotly.graph_objects as go

Z = hierarchy.linkage(df_categories, 'ward')

c, coph_dists = hierarchy.cophenet(Z, pdist(df_categories, 'hamming'))

ward = c

A = hierarchy.linkage(df_categories, 'average')

c, coph_dists = hierarchy.cophenet(A, pdist(df_categories, 'hamming'))

average = c

B = hierarchy.linkage(df_categories, 'single')

c, coph_dists = hierarchy.cophenet(B, pdist(df_categories, 'hamming'))

single = c

C = hierarchy.linkage(df_categories, 'complete')

c, coph_dists = hierarchy.cophenet(C, pdist(df_categories, 'hamming'))

complete = c

D = hierarchy.linkage(df_categories, 'weighted')

c, coph_dists = hierarchy.cophenet(D, pdist(df_categories, 'hamming'))

weighted = c

E = hierarchy.linkage(df_categories, 'centroid')

c, coph_dists = hierarchy.cophenet(E, pdist(df_categories, 'hamming'))

centroid = c

F = hierarchy.linkage(df_categories, 'median')

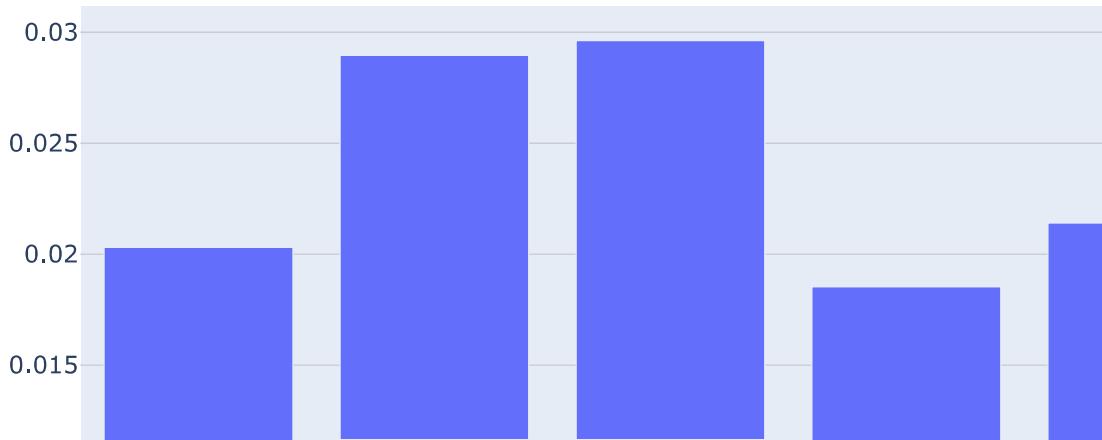
c, coph_dists = hierarchy.cophenet(F, pdist(df_categories, 'hamming'))

median = c

metrics=['ward', 'average', 'single', 'complete', 'weighted', 'centroid', 'median']

fig = go.Figure([go.Bar(x=metrics, y=[ward, average, single, complete, weighted, centroid,
```

```
fig.show()
```



In [80]:

```
from sklearn.metrics import silhouette_score
for n_clusters in range(2,10):
    clusterer = AgglomerativeClustering (n_clusters=n_clusters, distance_threshold = None)
    preds = clusterer.fit_predict(df_categories)

    score = silhouette_score (df_categories, preds)
    print ("For n_clusters = {}, silhouette score is {}".format(n_clusters, score))
```

```
For n_clusters = 2, silhouette score is 0.3418169738928389
For n_clusters = 3, silhouette score is 0.2545822953304492
For n_clusters = 4, silhouette score is 0.21919915815293894
For n_clusters = 5, silhouette score is 0.1857103264800873
For n_clusters = 6, silhouette score is 0.18638372714000376
For n_clusters = 7, silhouette score is 0.19230432552941282
For n_clusters = 8, silhouette score is 0.18314232248653736
For n_clusters = 9, silhouette score is 0.18527826016980006
```

In [81]:

```
#check the Davies-Bouldin score.We want a values as close to 0 as possible
from sklearn.metrics import davies_bouldin_score

for n_clusters in range(2,10):
    clusterer = AgglomerativeClustering (n_clusters=n_clusters, distance_threshold = None)
    preds = clusterer.fit_predict(df_categories)

    score = davies_bouldin_score (df_categories , preds)
    print ("For n_clusters = {}, the Davies-Bouldin score is {}".format(n_clusters, score))
```

For n\_clusters = 2, the Davies-Bouldin score is 1.1389848758478498  
 For n\_clusters = 3, the Davies-Bouldin score is 1.2779640449892402  
 For n\_clusters = 4, the Davies-Bouldin score is 1.3811925249494243  
 For n\_clusters = 5, the Davies-Bouldin score is 1.4498363376528345  
 For n\_clusters = 6, the Davies-Bouldin score is 1.4539406040366332  
 For n\_clusters = 7, the Davies-Bouldin score is 1.2906552438897336  
 For n\_clusters = 8, the Davies-Bouldin score is 1.2849398213510637  
 For n\_clusters = 9, the Davies-Bouldin score is 1.2524152683655723)

In [82]:

```
#Calinski-Harabasz index.We want as high a score as possible
from sklearn.metrics import calinski_harabasz_score

for n_clusters in range(2,10):
    clusterer = AgglomerativeClustering (n_clusters=n_clusters, distance_threshold = None)
    preds = clusterer.fit_predict(df_categories)

    score = calinski_harabasz_score(df_categories, preds)
    print ("For n_clusters = {}, the Calinski-Harabasz score is {}".format(n_clusters, sco
```

For n\_clusters = 2, the Calinski-Harabasz score is 519.8935666574885  
 For n\_clusters = 3, the Calinski-Harabasz score is 428.99352073446323  
 For n\_clusters = 4, the Calinski-Harabasz score is 368.76673452400416  
 For n\_clusters = 5, the Calinski-Harabasz score is 336.3946416980188  
 For n\_clusters = 6, the Calinski-Harabasz score is 315.7454361879021  
 For n\_clusters = 7, the Calinski-Harabasz score is 300.9252851740537  
 For n\_clusters = 8, the Calinski-Harabasz score is 291.83934867479667  
 For n\_clusters = 9, the Calinski-Harabasz score is 285.91597760644333)

In [83]:

```
model = AgglomerativeClustering(distance_threshold=0, n_clusters = None)
model = model.fit(df_categories)

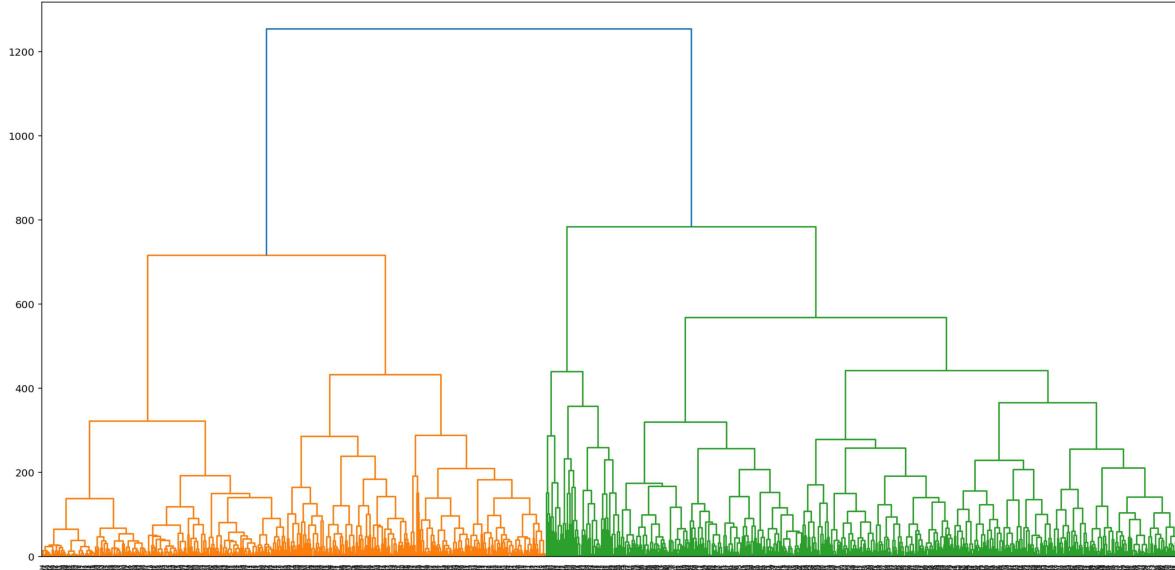
Z = hierarchy.linkage(model.children_, 'average')

plt.figure(figsize=(20,10))

dn = hierarchy.dendrogram(Z)
```

Out[83]:

&lt;Figure size 1440x720 with 0 Axes&gt;



In [84]:

```

df = pd.DataFrame(df_trvl_rvw)

model = AgglomerativeClustering(distance_threshold=None, n_clusters = 3)
model = model.fit(df_categories)
y_agg=model.fit_predict(df_categories)

df_agg = df_categories.copy()
df_agg["AggLabels"] = y_agg

#uncomment the following Line and let your machine explode!
sns.pairplot( df_agg, hue="AggLabels")

df_agg["AggLabels"].value_counts(0)

```

Out[84]:

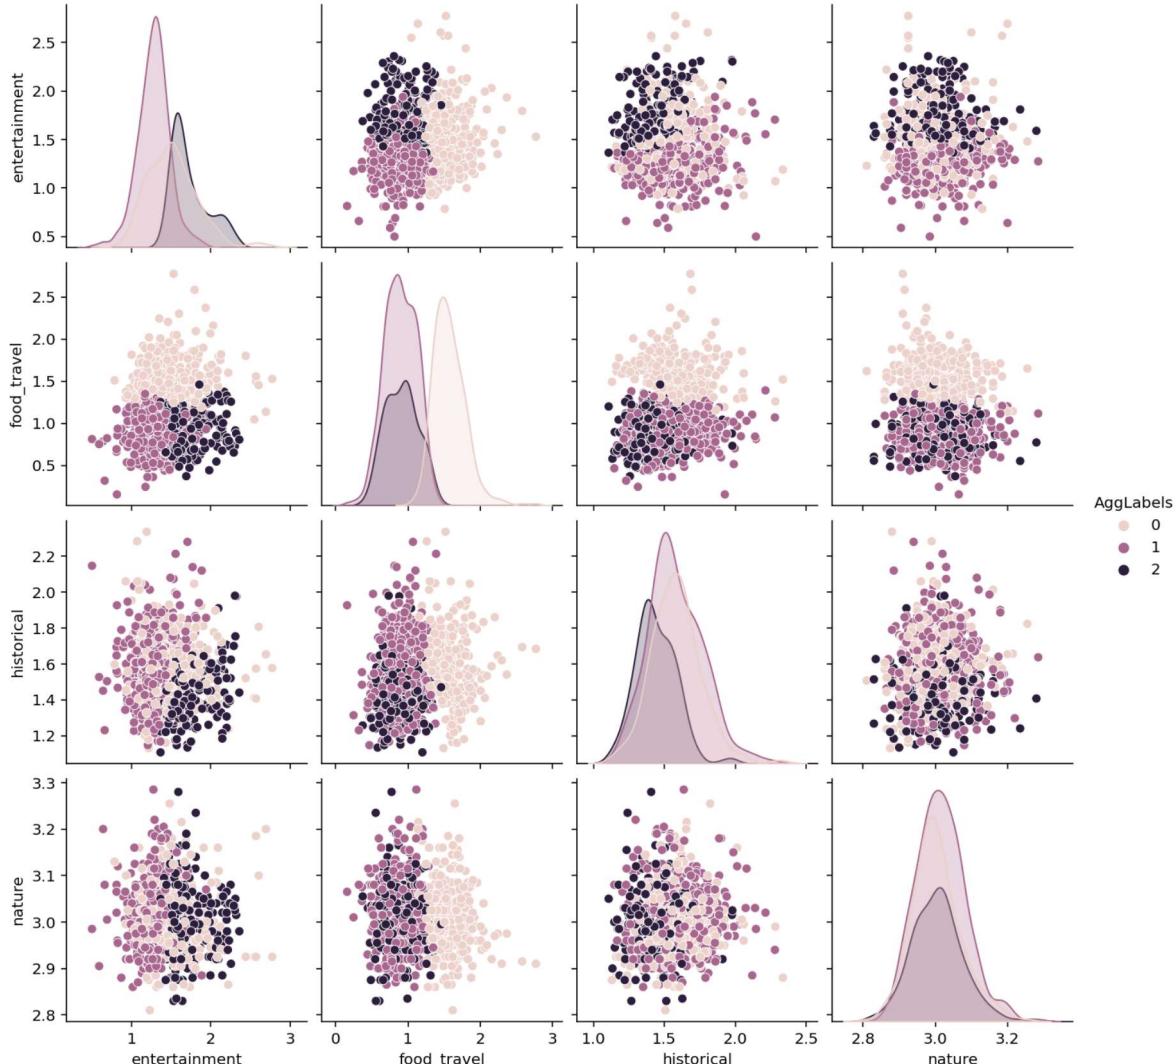
&lt;seaborn.axisgrid.PairGrid at 0x2284a80f7f0&gt;

Out[84]:

```

1    403
0    311
2    227
Name: AggLabels, dtype: int64

```



In [85]:

```
#DBSCAN
from sklearn.neighbors import NearestNeighbors

nearest_neighbors = NearestNeighbors(n_neighbors=73) #sqrt(5456) = 73
nearest_neighbors.fit(df_categories)
distances, indices = nearest_neighbors.kneighbors(df_categories)
distances = np.sort(distances, axis=0)[:, 1]
#print(distances)
plt.figure(figsize=(20,10))
plt.plot(distances)
plt.show()
```

Out[85]:

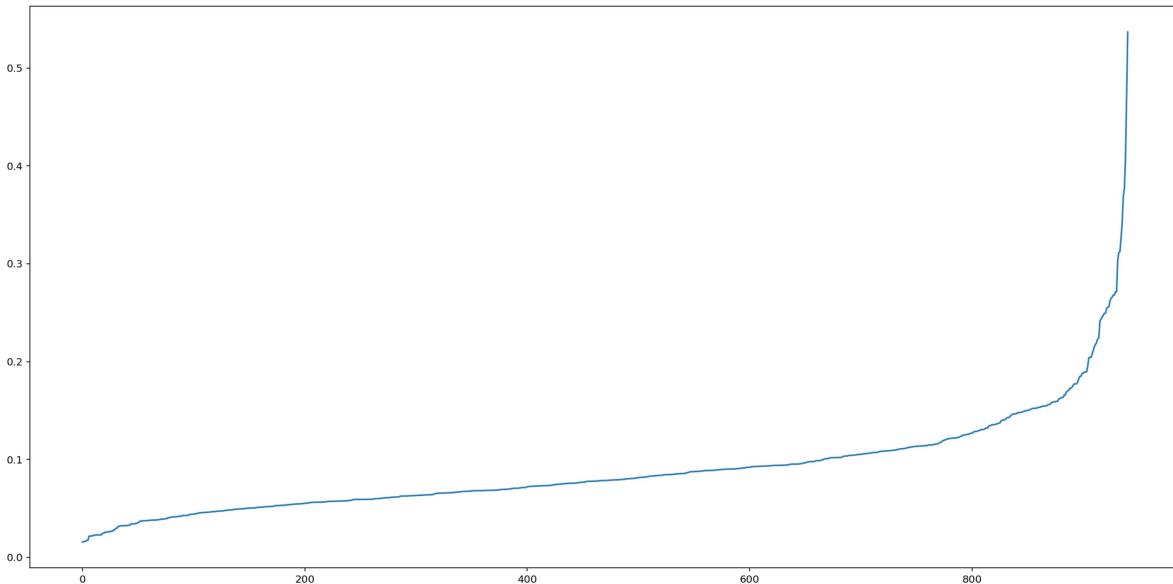
```
NearestNeighbors(n_neighbors=73)
```

Out[85]:

```
<Figure size 1440x720 with 0 Axes>
```

Out[85]:

```
[<matplotlib.lines.Line2D at 0x2285d323760>]
```



In [86]:

```
#Optimal value for eps where a 'knee' is formed is ~ 0.2
from sklearn.cluster import DBSCAN

model2 = DBSCAN(eps=0.3, min_samples=10)

model2 = model2.fit(df_categories)

np.unique(model2.labels_)

y_db=model2.labels_

df_db = df_categories.copy()
df_db["DBLabels"] = y_db

df_db[ "DBLabels" ].value_counts(0)

sns.pairplot( df_db, hue="DBLabels")
```

Out[86]:

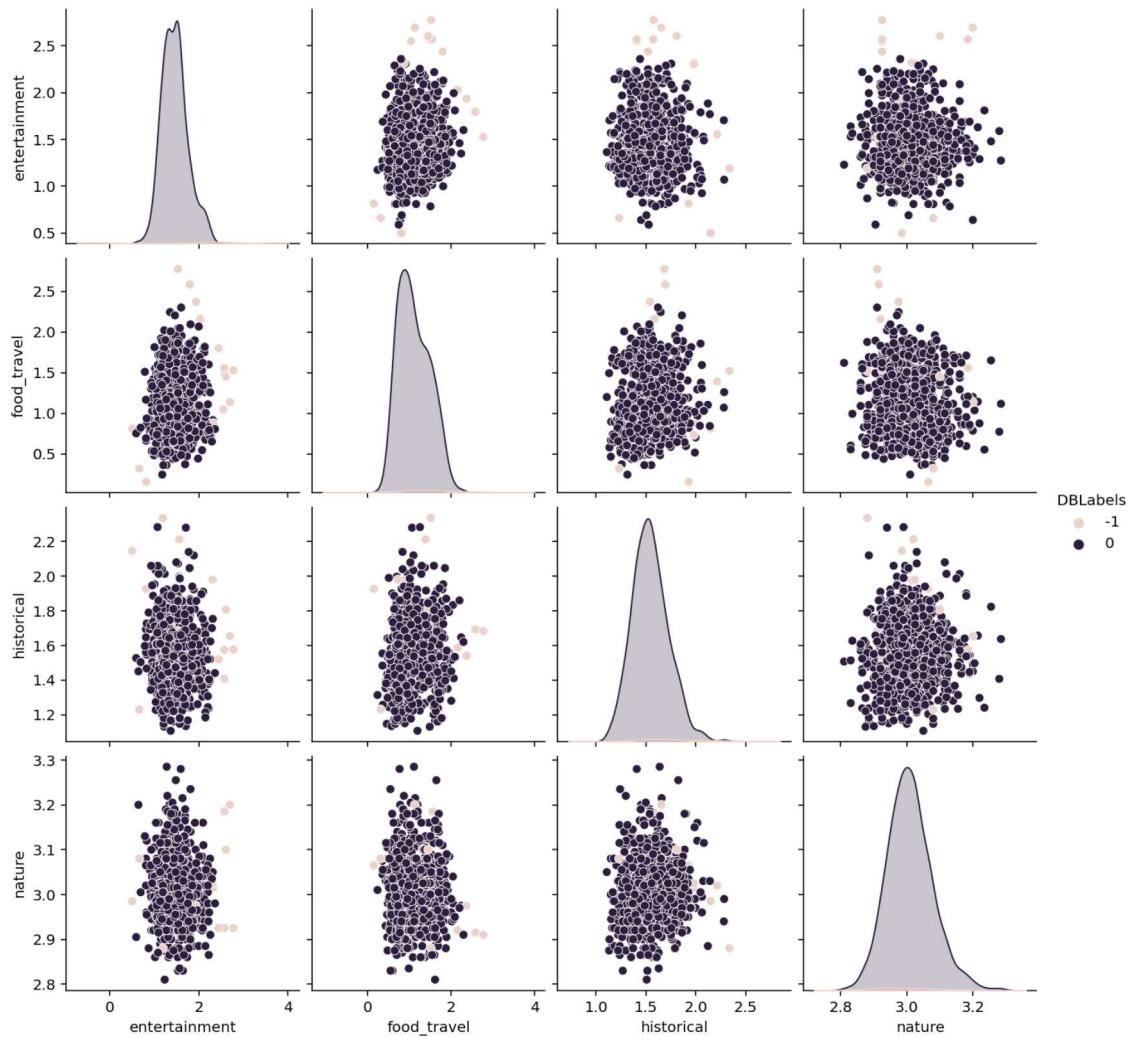
```
array([-1,  0], dtype=int64)
```

Out[86]:

```
0    923
-1    18
Name: DBLabels, dtype: int64
```

Out[86]:

```
<seaborn.axisgrid.PairGrid at 0x2285d3359a0>
```



In [87]:

```
df_db
```

Out[87]:

	entertainment	food_travel	historical	nature	DBLabels
0	1.810	1.776667	1.383333	2.990	0
1	2.030	2.160000	1.586667	2.920	-1
2	1.055	0.870000	1.320000	2.990	0
3	1.685	0.793333	1.256667	3.070	0
4	1.190	1.256667	1.530000	2.980	0
...	...	...	...	...	...
975	1.055	0.736667	1.606667	2.975	0
976	1.130	1.060000	1.610000	2.985	0
977	1.330	0.960000	1.643333	2.990	0
978	0.660	0.320000	1.230000	3.080	-1
979	0.950	1.333333	1.556667	3.025	0

941 rows × 5 columns

In [88]:

```
# Get names of indexes for which column DB_Labels has value -1
indexNames = df_db[ df_db['DBLabels'] == -1 ].index
# Delete these row indexes from DataFrame
df_db.drop(indexNames , inplace=True)

df_db['DBLabels'].unique()
```

Out[88]:

array([0], dtype=int64)

In [89]:

df\_categories

Out[89]:

	entertainment	food_travel	historical	nature
0	1.810	1.776667	1.383333	2.990
1	2.030	2.160000	1.586667	2.920
2	1.055	0.870000	1.320000	2.990
3	1.685	0.793333	1.256667	3.070
4	1.190	1.256667	1.530000	2.980
...	...	...	...	...
975	1.055	0.736667	1.606667	2.975
976	1.130	1.060000	1.610000	2.985
977	1.330	0.960000	1.643333	2.990
978	0.660	0.320000	1.230000	3.080
979	0.950	1.333333	1.556667	3.025

941 rows × 4 columns

In [90]:

```
import numpy as np
from sklearn.cluster import DBSCAN
from sklearn import metrics
from sklearn.datasets import make_blobs
from sklearn.preprocessing import StandardScaler
```

In [91]:

```
# Generate sample data
centers = df_categories
X, labels_true = make_blobs(n_samples=750, centers=df_categories, cluster_std=0.4,
                           random_state=0)

X = StandardScaler().fit_transform(X)
```

In [92]:

```
# Compute DBSCAN
db = DBSCAN(eps=0.3, min_samples=10).fit(X)
core_samples_mask = np.zeros_like(db.labels_, dtype=bool)
core_samples_mask[db.core_sample_indices_] = True
labels = db.labels_
```

In [93]:

```
# Number of clusters in labels, ignoring noise if present.
n_clusters_ = len(set(labels)) # - (1 if -1 in labels else 0)
n_noise_ = list(labels).count(-1)

print('Estimated number of clusters: %d' % n_clusters_)
print('Estimated number of noise points: %d' % n_noise_)
print("Homogeneity: %0.3f" % metrics.homogeneity_score(labels_true, labels))
print("Completeness: %0.3f" % metrics.completeness_score(labels_true, labels))
print("V-measure: %0.3f" % metrics.v_measure_score(labels_true, labels))
print("Adjusted Rand Index: %0.3f"
      % metrics.adjusted_rand_score(labels_true, labels))
print("Adjusted Mutual Information: %0.3f"
      % metrics.adjusted_mutual_info_score(labels_true, labels))
print("Silhouette Coefficient: %0.3f"
      % metrics.silhouette_score(X, labels))
```

Estimated number of clusters: 1  
 Estimated number of noise points: 750  
 Homogeneity: 0.000  
 Completeness: 1.000  
 V-measure: 0.000  
 Adjusted Rand Index: 0.000  
 Adjusted Mutual Information: 0.000

---

**ValueError** Traceback (most recent call last)  
 <ipython-input-93-7a73db259374> in <module>  
 13 % metrics.adjusted\_mutual\_info\_score(labels\_true, labels))  
 14 print("Silhouette Coefficient: %0.3f"  
--> 15 % metrics.silhouette\_score(X, labels))

c:\users\shery\pycharmprojects\cw\_ml\venv\lib\site-packages\sklearn\utils\validation.py in inner\_f(\*args, \*\*kwargs)  
 61 extra\_args = len(args) - len(all\_args)  
 62 if extra\_args <= 0:  
--> 63 return f(\*args, \*\*kwargs)  
 64  
 65 # extra\_args > 0

c:\users\shery\pycharmprojects\cw\_ml\venv\lib\site-packages\sklearn\metrics\cluster\\_unsupervised.py in silhouette\_score(X, labels, metric, sample\_size, random\_state, \*\*kwds)  
 115 else:  
 116 X, labels = X[indices], labels[indices]  
--> 117 return np.mean(silhouette\_samples(X, labels, metric=metric, \*\*kwds))  
 118  
 119

c:\users\shery\pycharmprojects\cw\_ml\venv\lib\site-packages\sklearn\utils\validation.py in inner\_f(\*args, \*\*kwargs)  
 61 extra\_args = len(args) - len(all\_args)  
 62 if extra\_args <= 0:  
--> 63 return f(\*args, \*\*kwargs)  
 64  
 65 # extra\_args > 0

c:\users\shery\pycharmprojects\cw\_ml\venv\lib\site-packages\sklearn\metrics

```
\cluster\_unsupervised.py in silhouette_samples(X, labels, metric, **kwds)
 227     n_samples = len(labels)
 228     label_freqs = np.bincount(labels)
--> 229     check_number_of_labels(len(le.classes_), n_samples)
 230
 231     kwds['metric'] = metric

c:\users\shery\pycharmprojects\cw_ml\venv\lib\site-packages\sklearn\metrics
\cluster\_unsupervised.py in check_number_of_labels(n_labels, n_samples)
 32     """
 33     if not 1 < n_labels < n_samples:
---> 34         raise ValueError("Number of labels is %d. Valid values are 2
"
 35                                         "to n_samples - 1 (inclusive)" % n_labels)
 36
```

**ValueError:** Number of labels is 1. Valid values are 2 to n\_samples - 1 (inclusive)

In [94]:

```
# Plot result
import matplotlib.pyplot as plt
%matplotlib inline
```

In [95]:

```
# Black removed and is used for noise instead.
unique_labels = set(labels)
colors = [plt.cm.Spectral(each)
          for each in np.linspace(0, 1, len(unique_labels))]
for k, col in zip(unique_labels, colors):
    if k == -1:
        # Black used for noise.
        col = [0, 0, 0, 1]

class_member_mask = (labels == k)

xy = X[class_member_mask & core_samples_mask]
plt.plot(xy[:, 0], xy[:, 1], 'o', markerfacecolor=tuple(col),
          markeredgecolor='k', markersize=14)

xy = X[class_member_mask & ~core_samples_mask]
plt.plot(xy[:, 0], xy[:, 1], 'o', markerfacecolor=tuple(col),
          markeredgecolor='k', markersize=6)

plt.title('Estimated number of clusters: %d' % n_clusters_)
plt.show()
```

Out[95]:

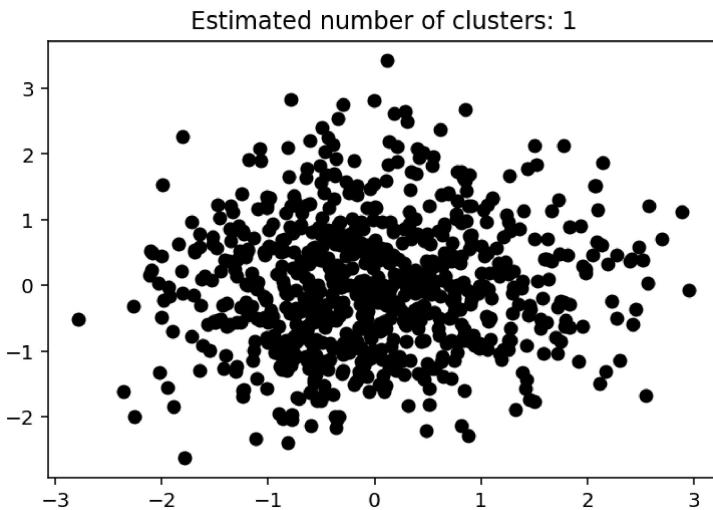
[&lt;matplotlib.lines.Line2D at 0x22859f28640&gt;]

Out[95]:

[&lt;matplotlib.lines.Line2D at 0x22859f288e0&gt;]

Out[95]:

Text(0.5, 1.0, 'Estimated number of clusters: 1')



In [97]:

```

import numpy as np
import matplotlib.pyplot as plt
from sklearn import metrics
from sklearn.datasets import make_circles
from sklearn.preprocessing import StandardScaler
from sklearn.cluster import DBSCAN
X, y = make_circles(n_samples=750, factor=0.3, noise=0.1)
X = StandardScaler().fit_transform(X)
y_pred = DBSCAN(eps=0.3, min_samples=10).fit_predict(X)

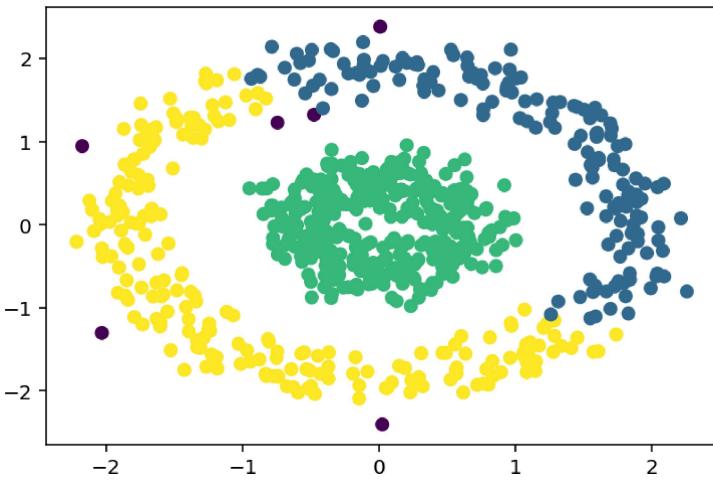
plt.scatter(X[:,0], X[:,1], c=y_pred)
print('Number of clusters: {}'.format(len(set(y_pred[np.where(y_pred != -1)]))))
print('Homogeneity: {}'.format(metrics.homogeneity_score(y, y_pred)))
print('Completeness: {}'.format(metrics.completeness_score(y, y_pred)))
print("V-measure: {:.3f} % metrics.v_measure_score(labels_true, y_pred))#Labels"))
print("Adjusted Rand Index: {:.3f}"
      % metrics.adjusted_rand_score(labels_true, y_pred))#Labels))
print("Adjusted Mutual Information: {:.3f}"
      % metrics.adjusted_mutual_info_score(labels_true, y_pred))#Labels))
print("Silhouette Coefficient: {:.3f}"
      % metrics.silhouette_score(X,y_pred))

```

Out[97]:

&lt;matplotlib.collections.PathCollection at 0x228591ed9a0&gt;

Number of clusters: 3  
 Homogeneity: 1.0  
 Completeness: 0.6477405612067773  
 V-measure: 0.278  
 Adjusted Rand Index: 0.000  
 Adjusted Mutual Information: 0.000  
 Silhouette Coefficient: 0.296



In [ ]: